

CPUNAME Instruction Set Architecture (ISA) Specification

lucchie-cmd

March 19, 2025

Contents

| | | |
|-----------|--|----------|
| 1 | General overview | 1 |
| 2 | Registers | 1 |
| 2.1 | NOTES | 1 |
| 3 | Instruction encoding | 2 |
| 3.1 | General Instruction Format | 2 |
| 4 | Instruction set | 3 |
| 4.1 | Instruction set as mnemonics | 3 |
| 4.1.1 | Moves | 3 |
| 4.1.2 | Stack Manipulation | 3 |
| 4.1.3 | Control flow | 4 |
| 4.1.4 | Arithmetic | 4 |
| 4.1.5 | Miscellaneous | 4 |
| 4.1.6 | Prefixes | 4 |
| 4.2 | Microinstructions | 5 |
| 4.2.1 | Loading | 5 |
| 4.2.2 | Writing | 5 |
| 4.2.3 | ALU | 5 |
| 4.2.4 | MMU | 6 |
| 5 | Privilege levels | 6 |
| 5.1 | Instructions | 6 |
| 5.2 | Registers | 6 |
| 6 | Memory | 6 |
| 7 | MMU | 7 |
| 8 | ALU | 8 |
| 9 | I/O | 8 |
| 9.1 | MMIO | 8 |
| 10 | Interrupts | 8 |

1 General overview

2 Registers

| Register | Purpose |
|----------|--|
| R1-R32 | General purpose registers |
| SP | Pointer to where the stack head currently is |
| IP | Pointer to where the CPU will execute instructions |
| PTP | Pointer to where the MMU will translate addresses |
| FLAGS | Status flags used for conditional jumps and moves |
| EADDR | A temporary effective address that the CPU uses to calculate some addresses for instructions |
| TEMP | A temporary register that the CPU uses for instructions that output nothing (Like COMP and TEST) |

Table 1: Register overview

See table above for register descriptions. Both SP and IP are virtual addresses meaning that they should be mapped in the PTP to be used. EADDR and TEMP registers are inaccessible for both the privileged and unprivileged. Flags and the PTP registers can only be set directly via the privileged program, otherwise they can only be affected by the ALU.

2.1 NOTES

The flags register affects conditional execution such as jumps and moves and can only be effected by the ALU

3 Instruction encoding

3.1 General Instruction Format

Unlike RISC, the CPUNAME will have variable length instructions, meaning that every instruction only takes up the binary space it needs. A typical instruction will generally look like so:

| Field | Size (In bytes) | Description |
|----------------------------------|------------------|---|
| Prefix | 1 | The instruction prefix, will change what addressing mode(s) to use. Described in Instruction Set subsection Prefixes |
| Opcode | 1 | The main instruction identifier. Contains number of sources and if the destination or source is a memory or register or immediate |
| Operands(s) (R8/16/32/64/128) | 1 | A register to use in the instruction |
| Operands(s) (M8/16/32/64/128) | 1, 2, 4, 8 or 16 | A memory location to use in the instruction. Number of bytes pulled from memory will depend on the instruction's prefix otherwise default 128 bit addressing is assumed |
| Operands(s) (I8/16/32/64/128) | 1, 2, 4, 8 or 16 | An immediate value to use in the instruction. Number of bytes will depend on the instruction's prefix otherwise default 128 bit addressing is assumed |

Table 2: Instruction overview

Prefixes and operands are optional, some instructions (Described in the instruction set table where the operands are None) don't use operands. An instruction that has a prefix byte but doesn't take any operands will trigger an Invalid Prefix Opcode interrupt (#IPO). The number of operands for a instruction if specified in the table in Instruction Set, please refer to that for specific instruction encoding and arguments.

4 Instruction set

4.1 Instruction set as mnemonics

4.1.1 Moves

| Bytes | Mnemonic | Operand 1 | Operand 2 | Description |
|-------|----------|-----------------|-------------------------------------|---|
| 00 | MOVE | R8/16/32/64/128 | R8/16/32/64/128 | Move value at source register (Operand 2) to destination register (Operand 1) |
| 01 | MOVE | R8/16/32/64/128 | M8/16/32/64/128 | Move value located at the memory address in (Operand 2) to destination register (Operand 1) |
| 02 | MOVE | R8/16/32/64/128 | R8/16/32/64/128+ I8/16/32/64/128 | Move value at the address in the register of Operand 2, plus the immediate value of Operand 2 into the destination register (Operand 1) |
| 03 | MOVE | R8/16/32/64/128 | R8/16/32/64/128+ R8/16/32/64/128 | Move value at the address in the register of Operand 2, plus the value of the register of Operand 2 into the destination register (Operand 1) |
| 04 | MOVE | M8/16/32/64/128 | R8/16/32/64/128 | Move the value of register Operand 2 into the memory location specified in Operand 1 |

4.1.2 Stack Manipulation

| Bytes | Mnemonic | Operand 1 | Operand 2 | Description |
|-------|----------|-----------------|-----------|---|
| 05 | PUSH | R8/16/32/64/128 | None | Push value in operand 1 on the stack and decrement stack pointer by N bytes (specified using prefix) |
| 06 | POP | R8/16/32/64/128 | None | Pop value at current stack pointer into Operand 1 and increment stack pointer by N bytes (specified using prefix) |

4.1.3 Control flow

| Bytes | Mnemonic | Operand 1 | Operand 2 | Description |
|-------|----------|-----------|-----------|--|
| 07 | JIE | I128 | None | Perform a conditional jump to the immediate value of Operand 1 if zero flag is set |
| 08 | JIL | I128 | None | Perform a jump to the immediate value of Operand 1 if SF != OF |
| 09 | JMP | I128 | None | Perform an unconditional jump to the immediate value of Operand 1 |
| 0A | CALL | I128 | None | Perform a call to the immediate value of Operand 1, push the value of the return address (Current IP+17) first |
| 0B | RET | None | None | Perform a return from a procedure by popping the return value from the stack |
| 0C | U2K | None | None | Performs a call to the elevated program to handle privileged instructions, changes the privilege level |

4.1.4 Arithmetic

| Bytes | Mnemonic | Operand 1 | Operand 2 | Description |
|-------|----------|-----------------|-----------------|--|
| 0D | TEST | R8/16/32/64/128 | R8/16/32/64/128 | Perform a bitwise AND but don't store the result only update the FLAGS |
| 0E | XOR | R8/16/32/64/128 | R8/16/32/64/128 | Perform a bitwise XOR On the operands and store the result in the first operand |
| 0F | SHL | R8/16/32/64/128 | I8/16/32/64/128 | Perform a shift to the left by the number of times specified in Operand 2 on Operand 1 |

4.1.5 Miscellaneous

| Bytes | Mnemonic | Operand 1 | Operand 2 | Description |
|-------|----------|-----------------|-----------------|---|
| 10 | COMP | R8/16/32/64/128 | I8/16/32/64/128 | Compare the value in register Operand 1 against the immediate value and set flags |

4.1.6 Prefixes

Certain prefixes can be placed before an instruction to modify its behavior. These prefixes can affect the selection of registers, immediate values, memory

addressing modes, or other instruction properties. By default, without any prefixes, all operands are 128 bits. The table below defines the available prefixes and their effects.

| Bytes | Effect |
|-------|---|
| 11 | Changes instruction to operate using 8 bit values for all operands |
| 12 | Changes instruction to operate using 16 bit values for all operands |
| 13 | Changes instruction to operate using 32 bit values for all operands |
| 14 | Changes instruction to operate using 64 bit values for all operands |

4.2 Microinstructions

4.2.1 Loading

| Mnemonic | Arguments | Description |
|----------|----------------|--|
| LDR | Register index | Load contents of the register specified in the register index onto the bus |
| LDI | Value | Load the provided value onto the bus |

4.2.2 Writing

| Mnemonic | Arguments | Description |
|----------|----------------|---|
| STR | Register index | Store the contents of the bus into the register specified in register index |

4.2.3 ALU

| Mnemonic | Arguments | Description |
|----------|-----------|--|
| PAA | None | Prepare ALU for an addition |
| LALU | Operand | Load the Nth ALU operand (with N being specified in the argument) from the bus |
| AE | None | Allow the ALU to output its result onto the bus |

4.2.4 MMU

| Mnemonic | Arguments | Description |
|----------|-----------|---|
| MR | None | Load value at the address specified in the EADDR register to the bus |
| MW | None | Write the value of the current bus to the address specified in the EADDR register |

5 Privilege levels

The CPUNAME and versions thereof have 2 privilege levels named unprivileged and privileged. These privilege levels can be compared to X86's ring 0 and ring 3 respectively. A program running in privileged mode can execute any and all instructions and is allowed to change things like the PTP register, SP and IP via direct assignment. However an unprivileged program running is not allowed to access some instructions and registers that a privileged program can access. For example the PTP register inaccessible for a unprivileged program and will throw an the Unprivileged Execute interrupt (`#UPE`). A full description of instructions and registers the unprivileged level isn't allowed to access will be provided in a list below.

5.1 Instructions

Currently there are no instructions that can only be run on a privileged level.

5.2 Registers

- PTP
- SP
- IP
- FLAGS
- EADDR
- TEMP

6 Memory

The CPUNAME can support up to 13 levels of paging, each taking up 9 bits of the virtual address with the exception of the last level which points to the physical address, this takes up 12 bits giving a total of 117 bits used for virtual addressing. The remaining 11 bits should all be the same, if for any reason these bits aren't the same when it is needed (Whether for reading or writing) it will trigger a Non Canonical interrupt (`#NCC`). The 12 bits are the default value and are extended by 9 bits for every paging level activated. The maximum

paging level useable is 13 by default but can be changed by setting a specific value in the paging control register (CRP) when the maximum paging level of 0 is applied, the MMU treats this as not having any paging enabled. For every new level of paging added, 9 bits of virtual address space get added which would equate to a function like: $2^{N \times 9}$ where N is the number of paging levels enabled.

The stack and instruction pointer (SP and IP respectively) are to be mapped in the PTP (If paging is enabled).

PTP describes a physical address stored somewhere in the RAM, the PTP is used for translating virtual addresses to physical ones via a table described below.

| Bits | Name | Description |
|------|------|---|
| 1 | RED | Describes if a page can be read from, trying to read from a non readable page will trigger an Un Readable Page interrupt (#URP). |
| 1 | PRS | Describes if a page is present, trying to use a non present page will trigger a Non Present Page interrupt (#NPP). |
| 1 | WRT | Describes if a page can be written to, trying to write to a non writeable page will trigger an Un Writeable Page interrupt (#UWP) |
| 1 | PVP | Describes if a page is occupied by the privileged program, trying to use a privileged page as an unprivileged program will trigger an Un Privilege Page interrupt (#UPP) |
| 1 | NEX | Describes if a page can *NOT* be executed, trying to execute from a non executable page results in a Non Executable Page interrupt (#NEP) |
| 40 | NLT | This describes the address of the next level of paging table, if this is the last level it contains the physical page number. Must be non null otherwise it will throw a Null Next Level interrupt (#NNL) |
| 19 | RSV | These bits are reserved for future purposes and should be zero at all times. Attempting to set these bits to any other value will result in a Non Zero Reserved interrupt (#NZR) |

7 MMU

This section describes how the CPUNAME's MMU (Memory Management Unit) works in detail

8 ALU

This section describes how the CPUNAME's ALU (Arithmetic Logic Unit) works in detail

9 I/O

The CPUNAME only support MMIO as the PIO isn't used anymore

9.1 MMIO

10 Interrupts

11 Pipeline